

# Operating system structures

## Chapter 2

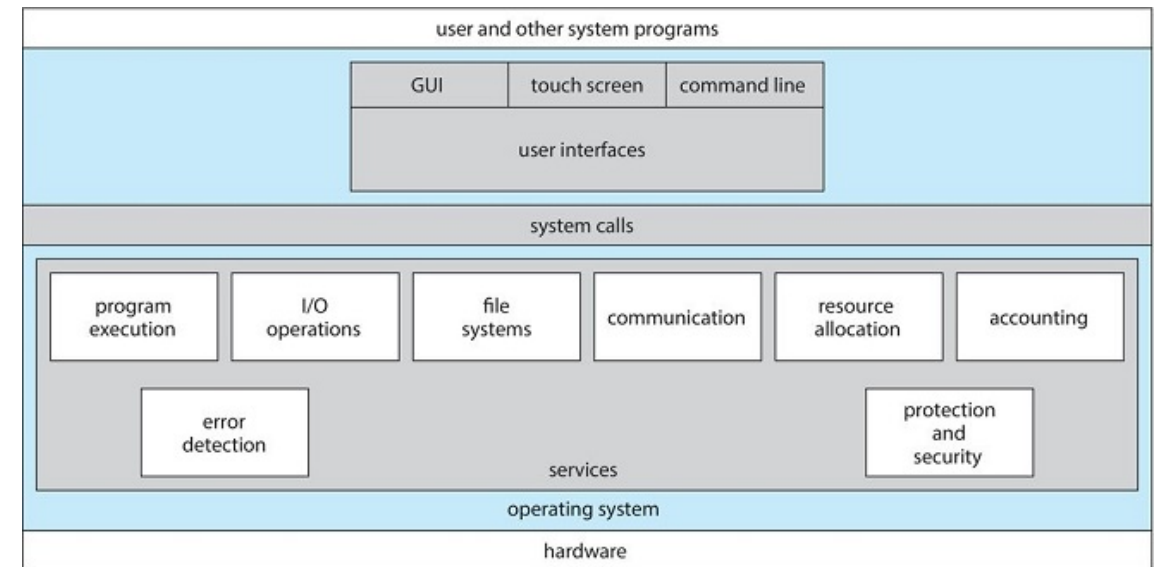


# 2.1 Operating system services: Multi-view perspective

## Internal structure vs external interface

- **View 1: services provided by system**
  - For applications and users
- **View 2: interface towards**
  - Users and programmers
  - Users: Gui or CLI (e.g.: cp f1.txt f2.txt)
  - Apps & programmers: API / system calls
- **View 3: components and interconnections of the OS**
  - Internal architecture/structure
  - Monolytics, layered, microkernel...

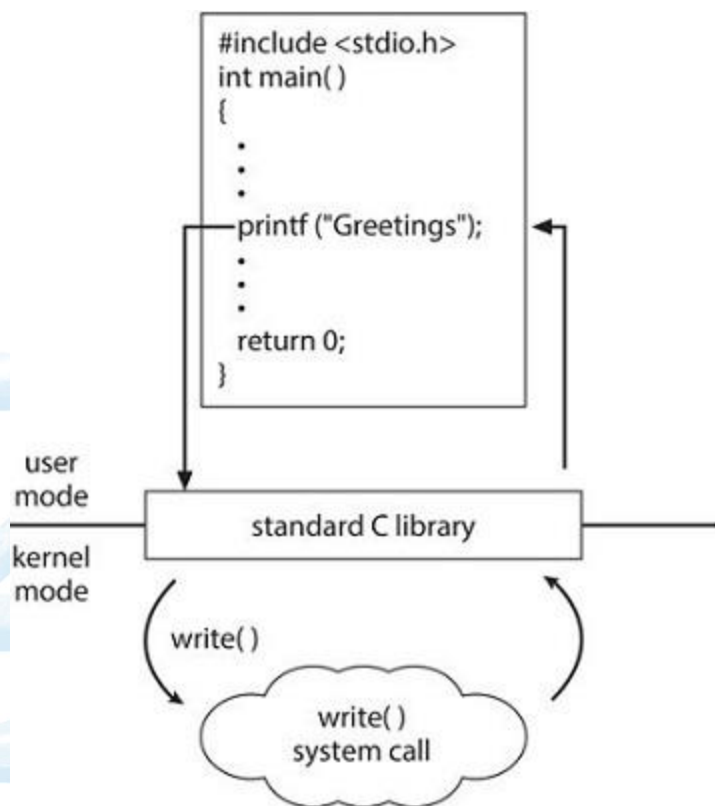
## Operating system services



# Example: printf("Greetings")

From API call to system call = 3 layers

## API call: printf() in stdio.h



## System call: write() → syscall()

- Write() is still a lower-level api call
- Syscall() is actual system call
- 1 = STDOUT\_FILENO

```
#include <unistd.h>
int main(void) {
    write(1, "hello, world!\n", 14);
    return 0;
}
```

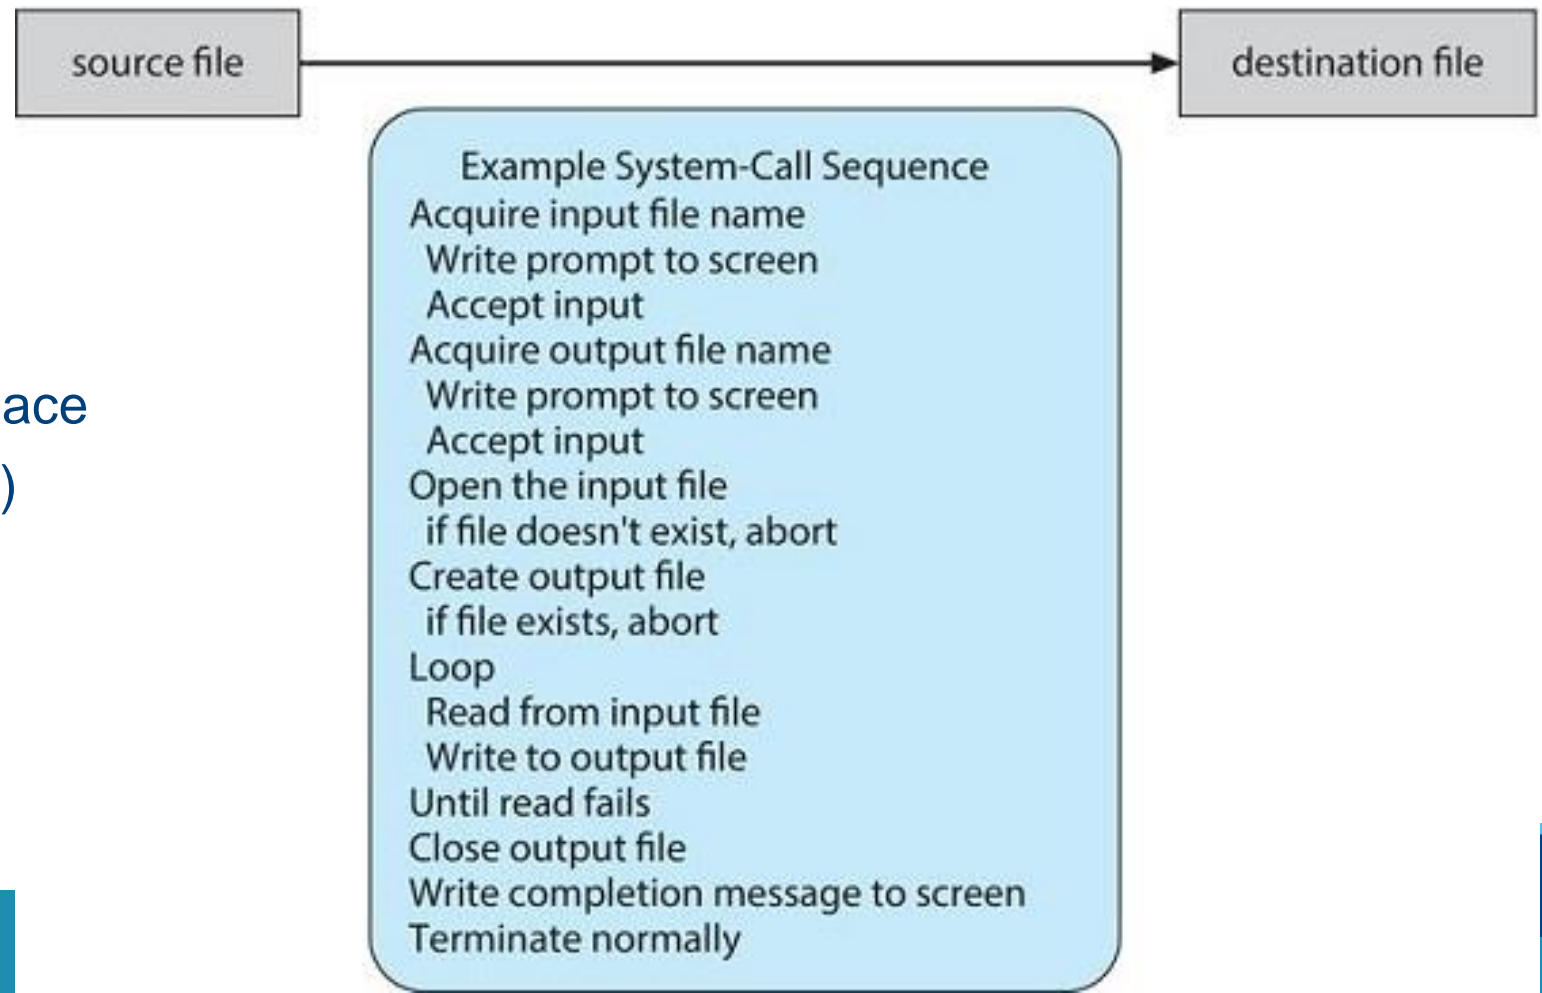
```
#include <unistd.h>
#include <sys/syscall.h>
int main(void) {
    syscall(SYS_write, 1, "hello, world!\n", 14);
    return 0;
}
```

## 2.3 System calls

### Interface

- To the OS services
- Functions
- Written in C / C++
- Some low-level tasks: ASM
- **Used by programs** in user space
  - OS programs/tools (e.g. cp)
  - Your applications
- Example
  - cp in.txt out.txt

### 2.3.1 E.g. “cp”: copy in.txt to out.txt



## 2.3.2 Application programming interface

### API

- **Even simple programs**
  - 1000s system calls per second
- **Too low level for application devs**
  - Solution: API
  - Higher-level functions for app devs
  - More stable, portable API
- **E.g.**
  - Win32 api for Windows
    - Abstracts over Win95/Win98/WinNT/Win10 kernel interface, which varies
  - POSIX api for linux/unix/macos
    - Kernel calls/interface abstracted

### Libc implements POSIX for C

- E.g. : read() and write()
- int fd: file descriptor id
- Void \*buf: buffer for the data
- Size\_t count: nb of bytes to read

```
#include <unistd.h>

ssize_t read(int fd, void *buf, size_t count)
```

Diagram illustrating the components of the `read()` function signature:

- `ssize_t`: return value
- `read`: function name
- `(int fd, void *buf, size_t count)`: parameters

## From API tot system call

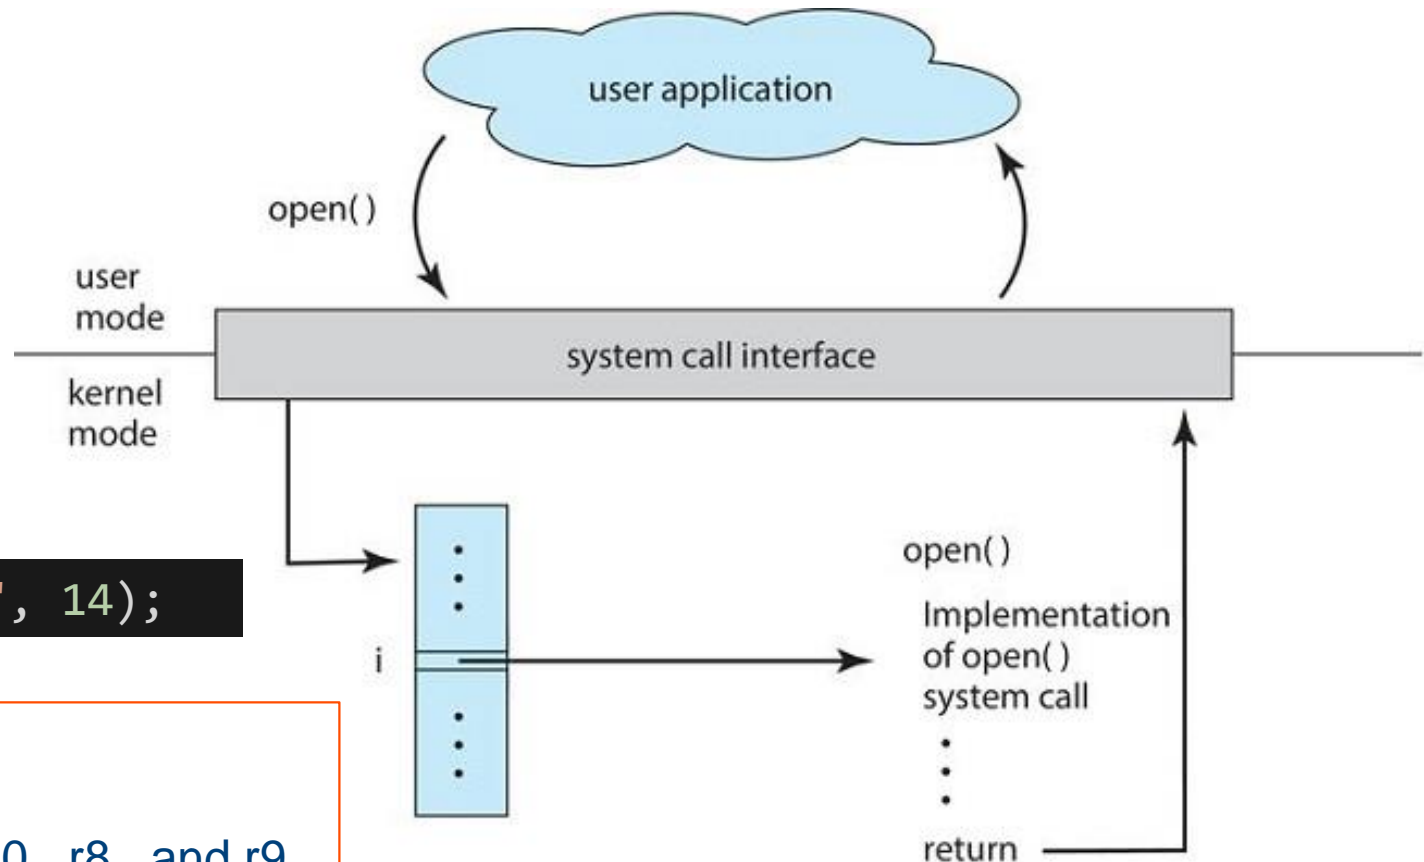
### From user mode to kernel mode

User application

- API
- System call
- Kernel mode
- Look up system call implementation
- Execute implementation
- Return

```
syscall(SYS_write, 1, "hello, world!\n", 14);
```

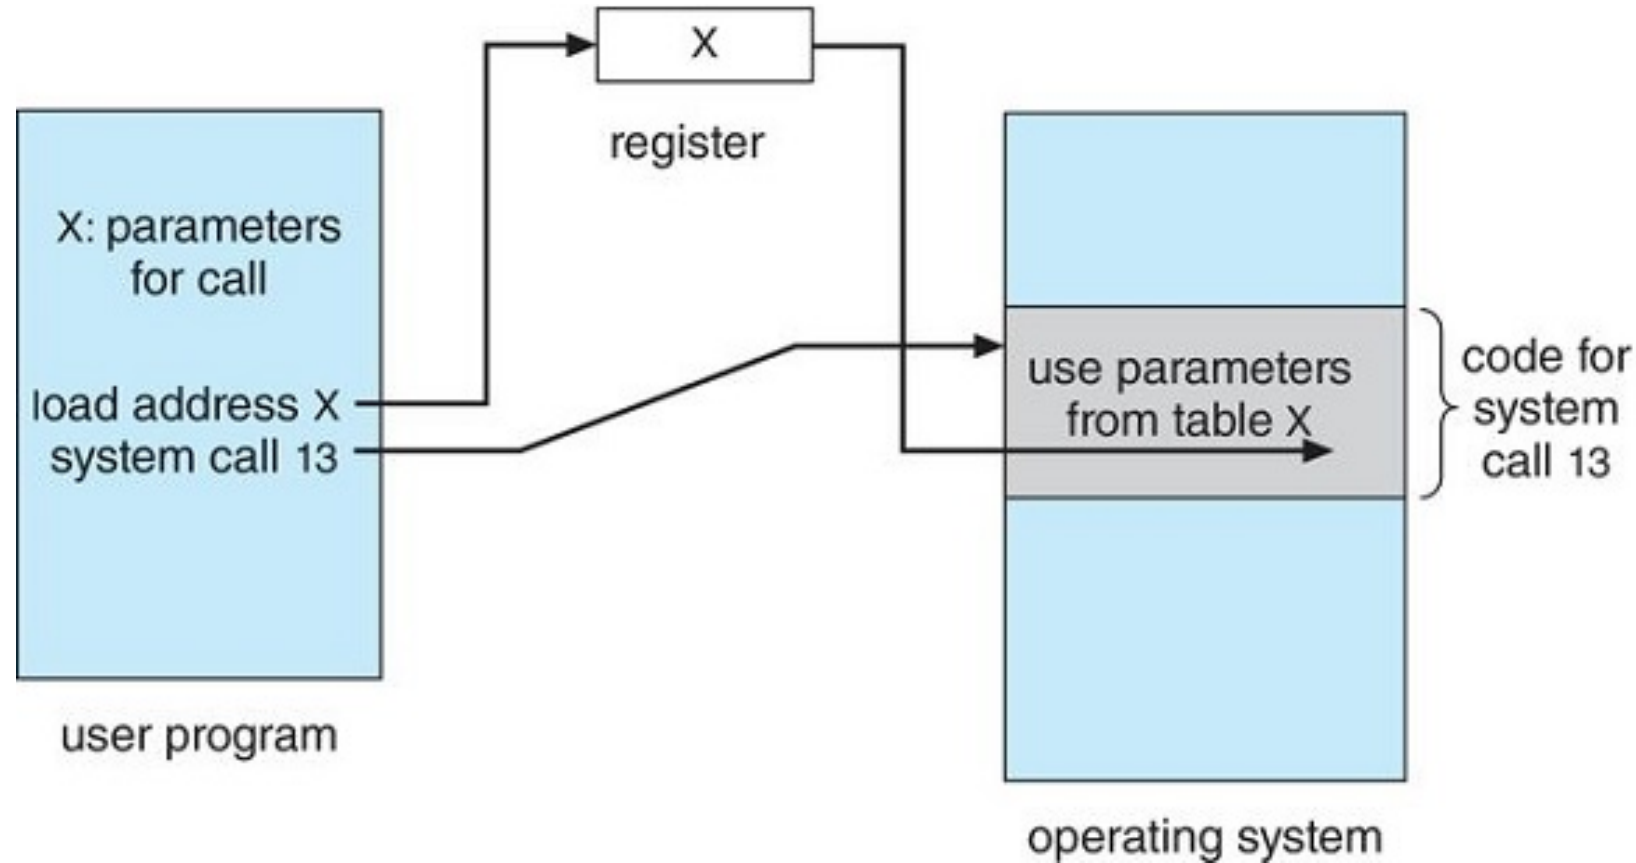
- To make a system call in 64-bit Linux,
- place the system call number in rax ,
  - then its arguments, in order, in rdi , rsi , rdx , r10 , r8 , and r9 ,
  - then invoke syscall instruction .
  - Some system calls return information, usually in rax .



## Passing parameters to OS

3 ways to pass params to system call

- via registers (for small set of params)
- Via block or table in memory
  - address via register
- User program pushes on stack,
  - popped by OS.



```
syscall(SYS_write, 1, "hello, world!\n", 14);
```



### 2.3.3 Types of system calls

#### Process control

- Create and terminate process
- Allocate and free memory

#### File management

- Create delete read write open close

#### Device management

- Attach device. Read/write.

#### Information management

- Time and date. timers

#### Communications

- Shared memory. Pipes.

#### Protection

- File permissions

### EXAMPLES OF WINDOWS AND UNIX SYSTEM CALLS

The following illustrates various equivalent system calls for Windows and UNIX operating systems.

	Windows	Unix
Process control	CreateProcess() ExitProcess() WaitForSingleObject()	fork() exit() wait()
File management	CreateFile() ReadFile() WriteFile() CloseHandle()	open() read() write() close()
Device management	SetConsoleMode() ReadConsole() WriteConsole()	ioctl() read() write()
Information maintenance	GetCurrentProcessID() SetTimer() Sleep()	getpid() alarm() sleep()
Communications	CreatePipe() CreateFileMapping() MapViewOfFile()	pipe() shm_open() mmap()
Protection	SetFileSecurity() InitializeSecurityDescriptor() SetSecurityDescriptorGroup()	chmod() umask() chown()



# 2.8 Operating system structure

Monoliths, modules and micro-kernels

## 2.8.1 Monolithic structure

“no structure at all”.

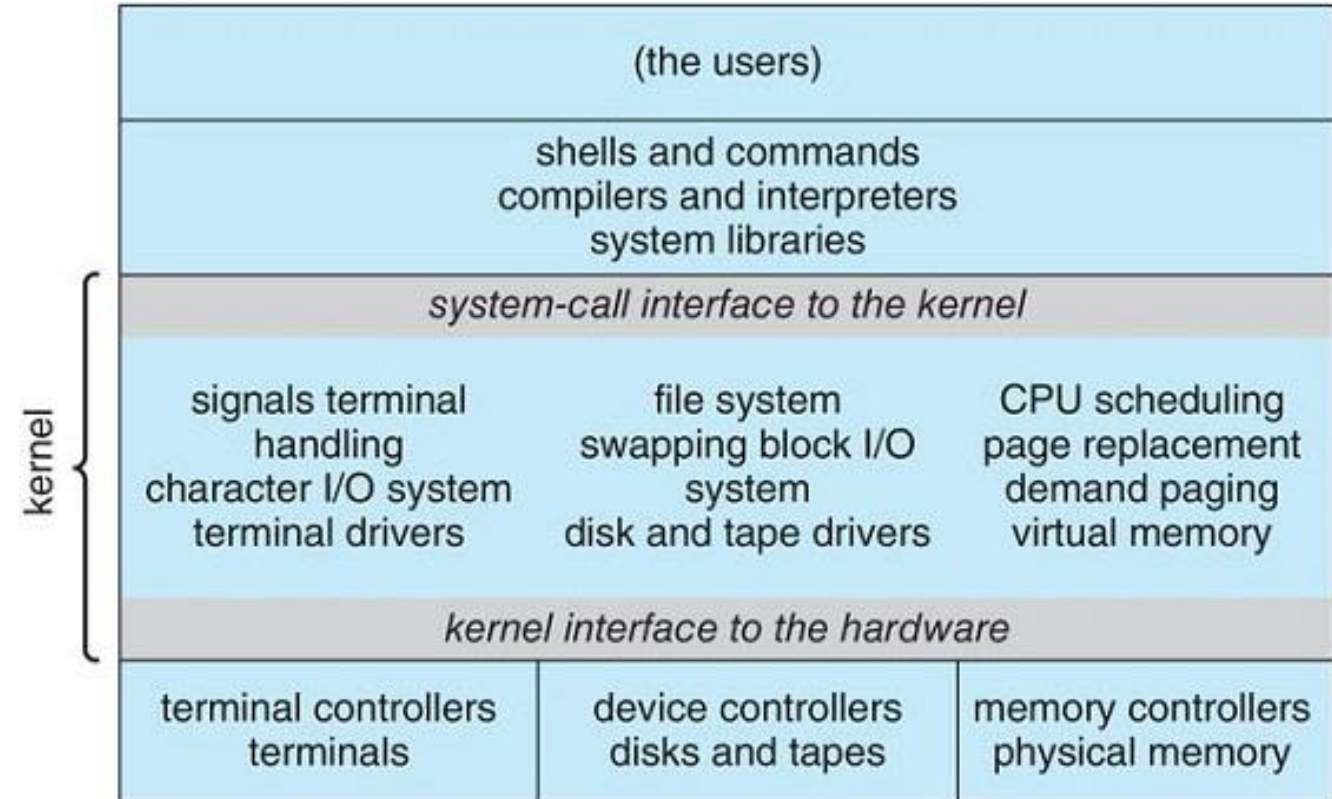
Kernel = 1 static binary file

e.g. UNIX

1 monolithic kernel + system programs

Kernel =

- Everything below system call interface
- Above physical hardware



## e.g. linux

Based on UNIX, structured similarly

Glibc = api for applications

Linux kernel = monolithic

- Runs entirely in a single address space
- Modular design for runtime modification

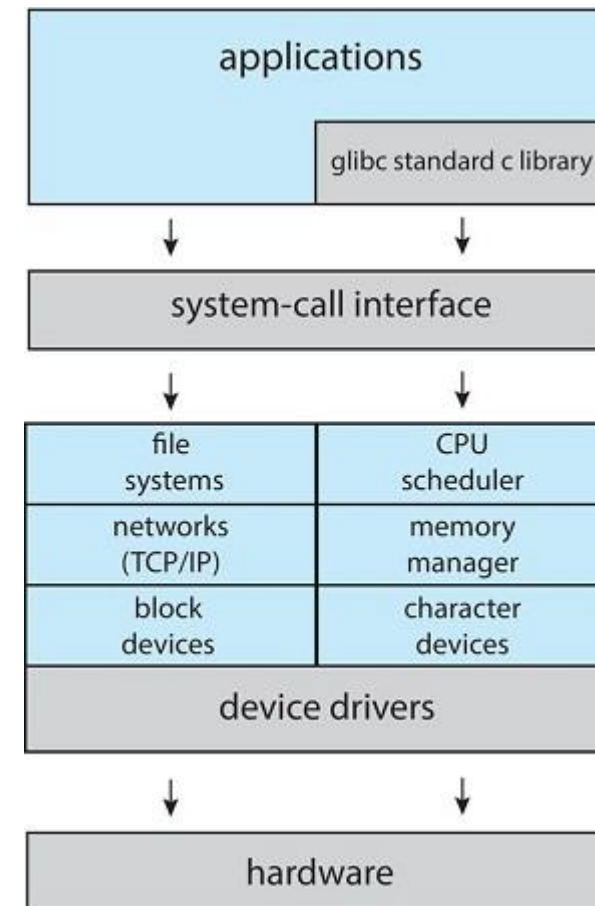
+ advantages

- performance, speed, efficiency

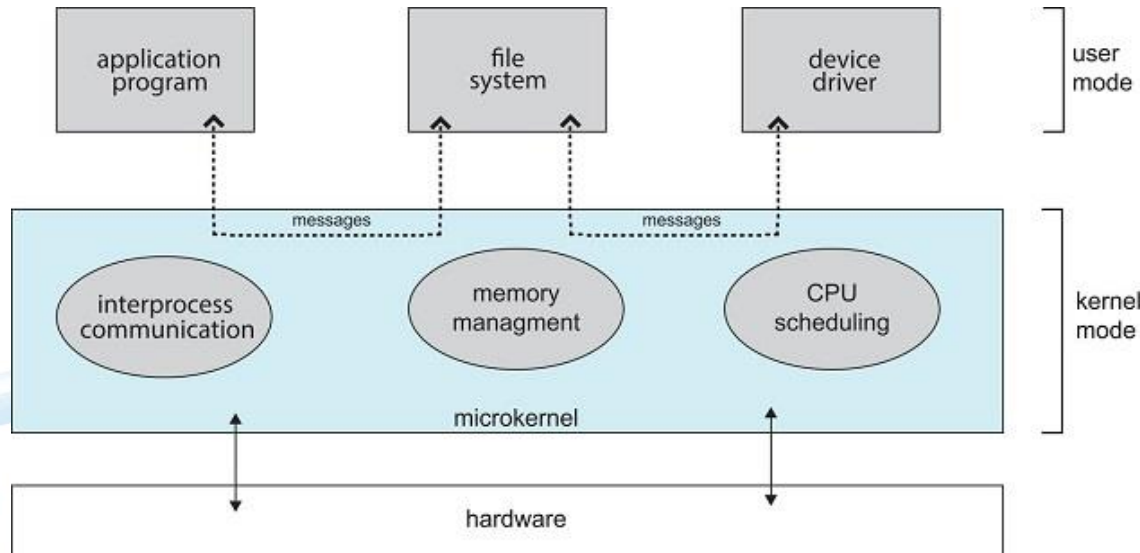
-disadvantages

Difficult to implement

Difficult to extend



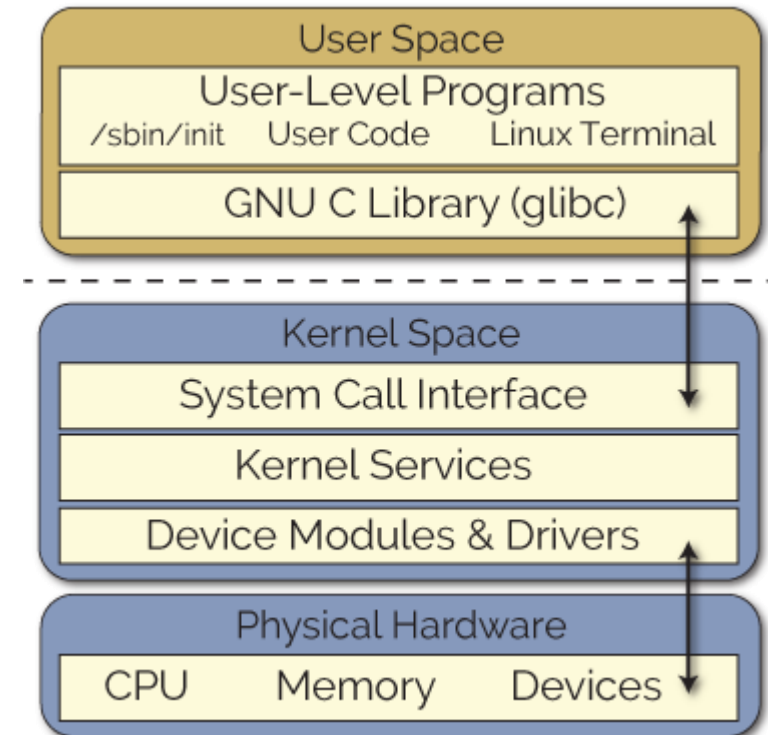
## 2.8.3 microkernels



- Modularization and minimization of the kernel
- e.g. Mach (1980s)
  - reaction against UNIX's unmanageability
- Removes all non-essential components from kernel
  - Implemented as user-level processes
  - Different address spaces
  - = much smaller kernel
- Today: Darwin kernel
  - Used by macOS and iOS.

## 2.8.4 Modules

- **Micro-kernel problems**
  - Have performance overhead
  - Lots of interprocess communication
  - Lots of context switches
- **Loadable kernel modules**
  - Run in kernel space
  - More efficient communication
  - Dynamically loaded
    - e.g. when usb drive plugged in
  - New services do not require kernel recompilation
    - E.g. for new file systems



# Quiz me quick


## **Microkernel = kernel**

- A) containing many components that are optimized to reduce resident memory size
- B) that is compressed before loading in order to reduce its resident memory size
- C) that is compiled to produce the smallest size possible when stored to disk
- D) that is stripped of all nonessential components

## **To load OS services dynamically you need**

- A) Virtual machines
- B) Modules
- C) File systems
- D) Graphical user interfaces





# Software engineering tools

gcc: compiling and linking.

Make, gdb, git.

## 2.5 Linkers and loaders

### Compilation (gcc)

- From source to **relocatable object file**
- Designed to be loaded into physical mem

### Linker (gcc)

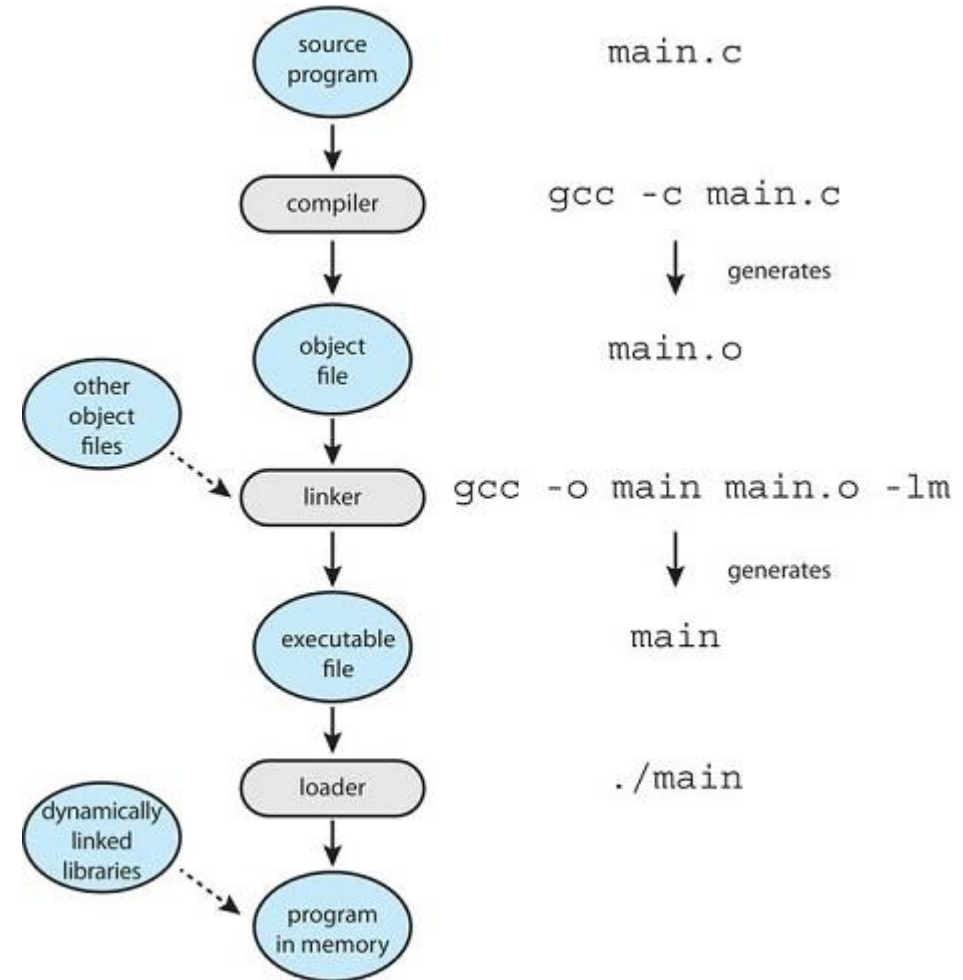
- Combines relocatable files into **executable**
- Includes other object files and libraries
- E.g. standard C lib (e.g. printf)
- E.g. math lib (-lm)
- UNIX: Executable and Linkable Format (ELF)

### Loader (./)

- Relocation to final address
- Adjusts code and data in program to match

### Dynamically linked library

- Loaded by need at **runtime**
- Shared by all programs that need it.



## Make: Complex compilation automation

Makefile:

```
hello:
    echo "hello world"
```

```
$ make
echo "hello world"
hello world
```

```
$ make blah
```

- Make is given blah as the target, so it first searches for this target
  - blah requires blah.o, so make searches for the blah.o target
  - blah.o requires blah.c, so make searches for the blah.c target
  - blah.c has no dependencies, so the echo command is run
- The gcc -c command is then run, because all of the blah.o dependencies are finished
- The top gcc command is run, because all the blah dependencies are finished
- That's it: blah is a compiled c program

```
blah: blah.o
gcc blah.o -o blah # Runs third
```

```
blah.o: blah.c
gcc -c blah.c -o blah.o # Runs second
```

```
blah.c:
echo "int main() { return 0; }" > blah.c # Runs first
```

## Gdb in a nutshell

### \$gdb -help

- **b** - Puts a breakpoint at the current line
- **b N** - Puts a breakpoint at line N
- **b fn** - Puts a breakpoint at the beginning of function "fn"
- **d N** - Deletes breakpoint number N
- **info break** - list breakpoints
- **r** - Runs the program until a breakpoint or error
- **c** - Continues running the program until the next breakpoint or error
- **f** - Runs until the current function is finished
- **s** - Runs the next line of the program
- **n** - Like s, but it does not step into functions
- **p var** - Prints the current value of the variable "var"
- **bt** - Prints a stack trace
- **u** - Goes up a level in the stack
- **d** - Goes down a level in the stack
- **q** - Quits gdb

## Adding debug info: gcc -g hello.c -o hello

### Run debugger: \$gdb hello

```
hello.c
1      #include <stdio.h>
2
3      int main(void)
4      {
5          printf("Hello, world!\n");
6
7          return 0;
8      }
9
10
11
12
13

child process 9054 In: main                               Line: 5      PC: 0x8048395
This GDB was configured as "i486-slackware-linux"...
(gdb) b main
Breakpoint 1 at 0x8048395: file hello.c, line 5.
(gdb) r
Starting program: /home/beej/hello

Breakpoint 1, main () at hello.c:5
(gdb) █
```

# Git version control

# get clone of remote repo on your local machine:

git clone <repo>

e.g:

\$ git clone git@host:username/repository.git.

\$ git clone [git@github.com:johndoe/my-app.git](https://github.com/johndoe/my-app.git).

# do work

→ You editing files in the repo.

# see changes

git status

git add main.c

git commit

#Push your changes to the remote repo on git server:

git push

